

COMP 7270/7276: “Advanced Algorithms”

Lecture 20: Randomized Algorithms

Bo Liu

Last Compiled: April 24, 2017

Probability Refresher

- Expectation of random variable:

$$\mathbb{E}[X] = \sum_r r \mathbb{P}[X = r]$$

- Linearity of expectation:

$$\mathbb{E}[X + Y] = \mathbb{E}[X] + \mathbb{E}[Y]$$

- Conditional Probability: For arbitrary events A and B ,

$$\mathbb{P}[A|B] = \mathbb{P}[A \cap B] / \mathbb{P}[B]$$

$$\text{and } \mathbb{P}[\cap_{i=1}^n A_i] = \mathbb{P}[A_1] \mathbb{P}[A_2|A_1] \dots \mathbb{P}[A_n | \cap_{i=1}^{n-1} A_i]$$

Outline

Randomized Algorithms

Quicksort

Karger's Randomized Min-Cut Algorithm

Deterministic Algorithms

- ▶ The algorithms we've seen so far have been deterministic.
- ▶ We want to aim for properties like
 - ▶ Good **worst-case** behavior.
 - ▶ Getting **exact** solutions.
- ▶ Much of our complexity arises from the fact that there is little flexibility here.
- ▶ Often find complex algorithms with nuanced correctness proofs.

Randomized Algorithms

- ▶ A **randomized algorithm** is an algorithm that incorporates randomness as part of its operation.
- ▶ Often aim for properties like
 - ▶ Good **average-case** behavior.
 - ▶ Getting exact answers **with high probability**.
 - ▶ Getting answers that are **close to the right** answer.
- ▶ Often find very simple algorithms with dense but clean analyses.

Randomized Algorithms

Two types of randomized algorithms:

- ▶ **Las Vegas** algorithms: the output is deterministic (and correct - it will always return a sorted list) but that the runtime is variable and affected by the randomization.
- ▶ **Monte Carlo** algorithms: The runtime of this algorithm is deterministic, but it is not guaranteed to always return the correct answer. Instead, it has a (hopefully small) probability of returning an incorrect answer.

Where We're Going

Motivating examples:

- ▶ Quicksort are **Las Vegas** algorithms: they always find the right answer, but might take a while to do so.
- ▶ Karger's algorithm is a **Monte Carlo** algorithm: it might not always find the right answer, but has dependable performance.

Outline

Randomized Algorithms

Quicksort

Karger's Randomized Min-Cut Algorithm

Quicksort

Problem: Sort an array of distinct values $X = [x_1, \dots, x_n]$

Algorithm

1. Pick a **pivot** $x \in X$ at random from the array
2. Construct new arrays $Y = [y_1, \dots, y_k]$, $Z = [z_1, \dots, z_{n-k-1}]$ where

$$y < x < z \text{ for all } y \in Y, z \in Z$$

3. Recursively sort Y and Z to get Y' and Z'
4. Return the array that concatenates Y' , x , and Z'

What's the expected number of comparisons performed in this algorithm?

Probability two items are compared

Lemma

Let a and b be the i -th and j -th smallest element of X where $i < j$.

$$\Pr[a \text{ is compared to } b] = \frac{2}{j - i + 1}$$

Probability two items are compared

Lemma

Let a and b be the i -th and j -th smallest element of X where $i < j$.

$$\Pr[a \text{ is compared to } b] = \frac{2}{j - i + 1}$$

Proof.

1. Consider $S = \{x \in X : a \leq x \leq b\}$



Probability two items are compared

Lemma

Let a and b be the i -th and j -th smallest element of X where $i < j$.

$$\Pr[a \text{ is compared to } b] = \frac{2}{j - i + 1}$$

Proof.

1. Consider $S = \{x \in X : a \leq x \leq b\}$
2. a and b are compared iff the first pivot chosen from S is either a or b



Probability two items are compared

Lemma

Let a and b be the i -th and j -th smallest element of X where $i < j$.

$$\Pr[a \text{ is compared to } b] = \frac{2}{j - i + 1}$$

Proof.

1. Consider $S = \{x \in X : a \leq x \leq b\}$
2. a and b are compared iff the first pivot chosen from S is either a or b
3. Elements of S are equally likely to be chosen as a pivot, so

$$\Pr[a \text{ is compared to } b] = \frac{2}{|S|} = \frac{2}{j - i + 1}$$



Expected Number of Comparisons

Lemma

Expected number of comparisons performs is $O(n \log n)$.

Proof.

1. Let $Z_{ij} = 1$ if the i -th smallest element is compared to j -th smallest element and $Z_{ij} = 0$ otherwise.
2. Number of comparisons: $\sum_{1 \leq i < j \leq n} Z_{ij}$
3. Expected number of comparisons:

$$\mathbb{E} \left[\sum_{1 \leq i < j \leq n} Z_{ij} \right] = \sum_{1 \leq i < j \leq n} \mathbb{E}[Z_{ij}] = \sum_{1 \leq i < j \leq n} \frac{2}{j - i + 1} = \sum_{j=2}^n \sum_{k=2}^j \frac{2}{k}$$

4. Because $H_n = 1 + 1/2 + 1/3 + \dots + 1/n = O(\log n)$,

$$\mathbb{E} \left[\sum_{1 \leq i < j \leq n} Z_{ij} \right] = O(n \log n)$$



Outline

Randomized Algorithms

Quicksort

Karger's Randomized Min-Cut Algorithm

Min-Cut Problem

Given an unweighted, multi-graph $G = (V, E)$, we want to partition V into V_1 and V_2 such that $|E \cap (V_1 \times V_2)|$ is minimized.

Algorithm

- ▶ *Contract a random edge $e = (u, v)$ and remove self-loops but not multi-edges*
- ▶ *Repeat until there are only 2 vertices remaining.*
- ▶ *Output the number of remaining edges.*

Let $|V| = n$ and $|E| = m$.

Karger's Algorithm


Given an edge (u, v) in a multigraph, we can **contract** u and v as follows:

- ▶ Delete all edges between u and v .
- ▶ Replace u and v with a new supernode uv .
- ▶ Replace all edges incident to u or v with edges incident to the supernode uv .

Karger's algorithm is as follows:


- ▶ If there are exactly two nodes left, stop. The edges crossing those nodes form a cut.
- ▶ Otherwise, pick a random edge, contract it, then repeat.

Karger's Algorithm




Step 1

Karger's Algorithm




Step 2

Karger's Algorithm




Step 3

Karger's Algorithm




Step 4

Karger's Algorithm



Step 5

Karger's Algorithm



Step 6

Karger's Algorithm

- ▶ Consider any cut $C = (S, V - S)$.
- ▶ If we never contract any edges crossing C , then Karger's algorithm will produce the cut C .
 - ▶ Initially, all nodes are in their own cluster.
 - ▶ Contracting an edge that does not cross the cut can only connect nodes that both belong to the same side of the cut.
 - ▶ Stops when two supernodes remain, which must be the sets S and $V - S$.

The Story So Far

- ▶ We now have the following: **Karger's algorithm produces cut C iff it never contracts an edge crossing C .**
- ▶ How does this relate to min cuts?
- ▶ Across all cuts, min cuts have **the lowest probability** of having an edge contracted.
- ▶ Fewer edges than all non-min cuts.
- ▶ Intuitively, we should be more likely to get a min cut than a non-min cut.
- ▶ What is the probability that we do get a min cut?

Correctness with low probability

Theorem

The algorithm is correct with probability at least

$$2/n^2$$

and never an underestimate.

Preliminaries

- ▶ Let $C = (V_1, V_2)$ be a specific minimum cut with $|C| = k$.
- ▶ What can you infer?

Preliminaries

- ▶ Let $C = (V_1, V_2)$ be a specific minimum cut with $|C| = k$.
- ▶ What can you infer?
 - ▶ Every vertex in G must have degree at least k .
 - ▶ Since there are n vertices in G , there are at least $nk/2$ edges in G .
- ▶ $\Pr[\text{first edge chosen randomly} \in C] \leq \frac{k}{nk/2} = \frac{2}{n}$.

Preliminaries

- ▶ Suppose at some step of the algorithm, there are Graph G_ℓ with ℓ vertices left.
- ▶ What can you infer?

Preliminaries

- ▶ Suppose at some step of the algorithm, there are Graph G_ℓ with ℓ vertices left.
- ▶ What can you infer?
 - ▶ the size of the minimum cut in G_ℓ is at least k .
 - ▶ there are at least $\ell k/2$ edges in G_ℓ .
- ▶ $\Pr[\text{C is hit when there are } \ell \text{ vertices left} \mid \text{C is not hit before}]$
 $\leq \frac{k}{\ell k/2} = \frac{2}{\ell}$

Correctness with low probability

Proof.

- ▶ Minimum cut of the graph doesn't decrease.
- ▶ Let $C = (V_1, V_2)$ be a specific minimum cut with $|C| = k$.
- ▶ Let A_i be event that we don't contract edge across C at step i .

$$\mathbb{P}[\cap_{1 \leq i \leq n-2} A_i] = \mathbb{P}[A_1] \mathbb{P}[A_2|A_1] \dots \mathbb{P}[A_{n-2} | \cap_{1 \leq i \leq n-3} A_i]$$

- ▶ Number of edges before i -th step if no edges across C have been contracted so far is at least $(n-i+1)k/2$
- ▶ $\mathbb{P}[A_i | A_1 \cap A_2 \cap \dots \cap A_{i-1}] \geq 1 - 2/(n-i+1)$ and so

$$\begin{aligned} \mathbb{P}[\cap_{1 \leq i \leq n-2} A_i] &\geq (1 - \frac{2}{n})(1 - \frac{2}{n-1})(1 - \frac{2}{n-2}) \dots (1 - \frac{2}{3}) \\ &= \frac{n-2}{n} \cdot \frac{n-3}{n-1} \cdot \frac{n-4}{n-2} \cdot \dots \cdot \frac{1}{3} = \frac{2}{n(n-1)} \end{aligned}$$



Content

We will see

- ▶ Boosting The Probability
- ▶ Speeding Things Up

Boosting the probability

Theorem

Repeating $\alpha n^2/2$ times (with new random coin flips) and returning smallest cut is correct with probability at least $1 - e^{-\alpha}$.

Proof.



Boosting the probability

Theorem

Repeating $\alpha n^2/2$ times (with new random coin flips) and returning smallest cut is correct with probability at least $1 - e^{-\alpha}$.

Proof.

- ▶ Because each repeat is independent,

$$\mathbb{P}[\text{always fails}] = \prod_{1 \leq i \leq \alpha n^2/2} \mathbb{P}[i\text{-th try fails}] \leq (1 - 2/n^2)^{\alpha n^2/2}$$

- ▶ Use fact $1 - x \leq e^{-x}$ for $x \geq 0$ and simplify.



Speeding Things Up: The Karger-Stein Algorithm

Some Quick History

- ▶ David Karger developed the contraction algorithm in 1993. Its runtime was $O(n^4 \log n)$.
- ▶ In 1996, David Karger and Clifford Stein (the S in CLRS) published an improved version of the algorithm that is *dramatically* faster.
- ▶ **The Good News:** The algorithm makes intuitive sense.
- ▶ **The Bad News:** Some of the math is really, really hard.

Some Observations

- ▶ Karger's algorithm only fails if it contracts an edge in the min cut.
- ▶ The probability of contracting the wrong edge increases as the number of supernodes decreases.
- ▶ Since failures are more likely later in the algorithm, repeat just the later stages of the algorithm when the algorithm fails.

Intelligent Restarts

- ▶ Interesting fact: If we contract from n nodes down to $n/\sqrt{2}$ nodes, the probability that we don't contract an edge in the min cut is about 50%.
 - ▶ Can work out the math yourself if you'd like.
- ▶ What happens if we do the following?
 - ▶ Contract down to $n/\sqrt{2}$ nodes.
 - ▶ Run two passes of the contraction algorithm from this point.
 - ▶ Take the better of the two cuts.

The Success Probability

- ▶ This algorithm finds a min cut iff
 - ▶ The partial contraction step doesn't contract an edge in the min cut, and
 - ▶ At least one of the two remaining contractions does find a min cut.
- ▶ The first step succeeds with probability around 50%.
- ▶ Each remaining call succeeds with probability at least $4/n(n-1)$.
Thinking assignment to calculate why.

A Success Story

- ▶ This new algorithm has roughly twice the success probability as the original algorithm!
- ▶ Key Insight: Keep repeating this process!
 - ▶ Base case: When size is some small constant, just brute-force the answer.
 - ▶ Otherwise, contract down to $n/\sqrt{2}$ nodes, then recursively apply this algorithm twice to the remaining graph and take the better of the two results.
- ▶ This is the **Karger-Stein** algorithm.

Two Questions

- ▶ What is the success probability of this new algorithm?
 - ▶ This is extremely difficult to determine.
 - ▶ We'll talk about it later.
- ▶ What is the runtime of this new algorithm.
 - ▶ Let's use the Master Theorem.

The Runtime

- ▶ We have the following recurrence relation:

$$\begin{array}{ll} T(n) = c & \text{if } n \leq n_0 \\ T(n) = 2T(n/\sqrt{2}) + O(n^2) & \text{otherwise} \end{array}$$

- ▶ What does the Master Theorem say about it?

$$T(n) = O(n^2 \log n)$$

The Accuracy

- ▶ By solving a very tricky recurrence relation, we can show that this algorithm returns a min cut with probability $\Omega(1/\log n)$.
- ▶ If we run the algorithm roughly $\ln^2 n$ times, the probability that *all* runs fail is roughly

$$\left(1 - \frac{1}{\ln n}\right)^{\ln^2 n} \leq \left(\frac{1}{e}\right)^{\ln n} = \frac{1}{n}$$

- ▶ **Theorem:** The Karger-Stein algorithm is an $O(n^2 \log^3 n)$ -time algorithm for finding a min cut with high probability.

Major Ideas

- ▶ You can increase the success rate of a Monte Carlo algorithm by iterating it multiple times and taking the best option found.
- ▶ If you're more intelligent about how you iterate the algorithm, you can often do much better than this.